

Scaling PostgreSQL:

The Bad, The Good And The Ugly

London PostgreSQL Meetup 2019

Chris Ellis - @intrbiz

Hello!

- I'm Chris
 - IT jack of all trades
- Been using PostgreSQL for about 14 years
- Very much into Open Source
 - Started Bergamot Monitoring - open distributed monitoring
- Worked on various PostgreSQL systems
 - IoT TV Set top boxes
 - Smart energy meter analytics
 - Mixes of OLTP and OLAP workloads
 - Scaled PostgreSQL in various ways for various situations

The Bad: Scaling is hard

- Computers are still governed by the laws of physics
 - Speed of light is fixed
- Scaling databases is hard
 - There are no magic solutions
 - How to scale is very dependent upon your workload
 - Scaling for readers vs writers is very different
 - Scaling OLAP vs OLTP
 - Monolithic vs Micro-Service architectures
- Thankfully PostgreSQL is pretty flexible
- Your options might be limited by where your databases exist
 - Cloud vs on-prem
 - Managed vs unmanaged

The Good: Optimising


- You should start by utilising your resources as effectively as possible
 - Tune and optimise PostgreSQL for your use cases
 - Ensure things stay good - statistics, etc
- PostgreSQL has a huge range of features to help you
 - Indexes
 - There are more than just B-Trees
 - Expression indexes
 - Parallelisation
 - PostgreSQL has increasing support for executing queries across many CPUs
 - Partitioning
 - Handling large volumes of data can become problematic

The Good: A Bigger Boat

- Sometimes you're going to need a bigger boat
- Sometimes the easiest way to scale is to buy a bigger bit of hardware
 - Don't underestimate how effective this can be
 - Probably not as expensive as you think, compared to other options
- This is very dependent upon your workload
- If you're moving to a bigger box, you're going to need to change your ``postgresql.conf``
- Different workloads can be scaled up in different ways
 - You might just need fewer faster CPUs
 - You might just need a box of disks


The Good: Replication

- For workloads which are majority read
- Streaming replication provides a good simple way to replicate your data
- Backend systems can still interact with your single primary and read workload can run from replicas
- Proxies/Poolers such as pgBouncer, HAProxy or in app can redirect writes



The Good: Logical Replication

- For more complex workloads which are more a mix of read / write
- Logical replication can be used to split data across different servers
- Useful when there might be multiple backend systems responsible for different datasets but you want to present them together




The Ugly: Scale Out

- Sometime to scale, you need to split your workload across multiple servers
 - This works well for write heavy and complex workloads
- This is where life starts to get very complex and very dependent upon your workload.
 - And also where the most snake-oil seems to exist
- You need to carefully think about how you split and place your data
 - Even modern fast networks are very latent compared to memory or local disk
 - You need to think about co-location of data and denormalisation
 - You need to accept that some queries will get faster, some slower

The Ugly: Scale Out - FDW

- PostgreSQL has had Foreign Data Wrappers for a while now
 - Allows PostgreSQL to access data held in foreign systems
 - They even support writing to foreign systems
- Allows you to build some interesting sharded systems
 - Write directly into the shards or via the primary
 - Foreign tables can be used in inheritance, allowing you to query across all shards
 - Predicate push down can help select shard to access data from



The Ugly: Scale Out - PL/Proxy

- PL/Proxy is probably the oldest way to scale out PostgreSQL
 - Created by Skype, has a very particular way of doing things
 - It is by no means a magic solution
- If you love functions, it offers a huge amount of power and flexibility in how to access and store data across a cluster of databases
- Implemented as a procedural language
 - You create proxy functions which know how to call functions on remote shards
 - Gives control over how to partition data very specifically

The Ugly: Scale Out - PL/Proxy

```
CREATE OR REPLACE FUNCTION stb.stb_upsert(i_id BIGINT, i_mac_address MACADDR,  
...)
```

```
RETURNS BOOLEAN
```

```
LANGUAGE PLPROXY AS $$
```

```
    CLUSTER stb._get_cluster(true);
```

```
    RUN ON util.get_shard(i_id);
```

```
$$;
```

The Ugly: Scale Out - PL/Proxy

```
CREATE OR REPLACE FUNCTION stb.stb_count()  
  
RETURNS SETOF stb.t_table_row_count  
  
LANGUAGE PLPROXY AS $stb_count$  
  
    CLUSTER stb._get_cluster(FALSE);  
  
    RUN ON ALL;  
  
    SELECT inet_server_addr() AS "server_inet", current_database() AS "cdb",  
count(*) AS "cnt" FROM stb.stb;  
  
$stb_count$;
```

I could go on, but won't

- This has been a short overview of some methods
 - Reality will be you'll need to combine quite a few of these
- There are many other projects out there:
 - BDR, Citus, PostgreSQL-XL, Timescale DB to name but a few
- Questions?